

# Neural Networks I

## Lecture 18

# What's the hype around neural networks?

Character/handwriting recognition

Image compression (autoencoders)






Stock market prediction

Credit approval

And some very recent interesting deep learning applications...



# Image Style Transfer



Dumoulin, Vincent, Jonathon Shlens, and Manjunath Kudlur. "A learned representation for artistic style." CoRR, abs/1610.07629 2.4 (2016): 5.

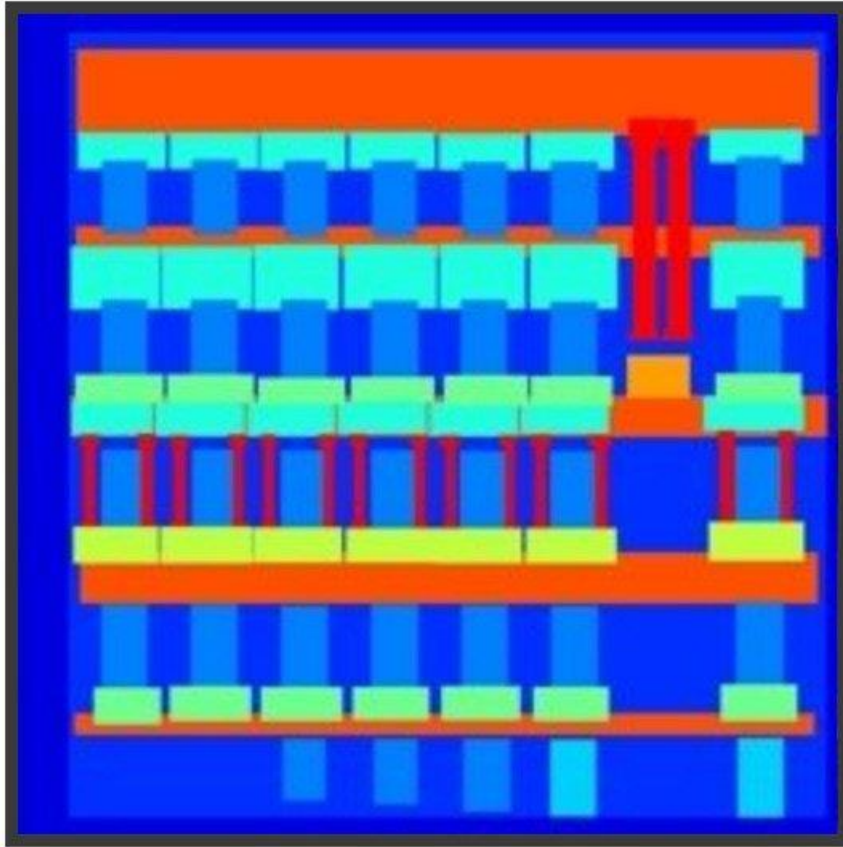


# Image-to-image translation

TOOL

- background
- wall
- door
- window**
- window sill
- window head
- shutter
- balcony
- trim
- cornice
- column
- entrance

INPUT



pix2pix  
process

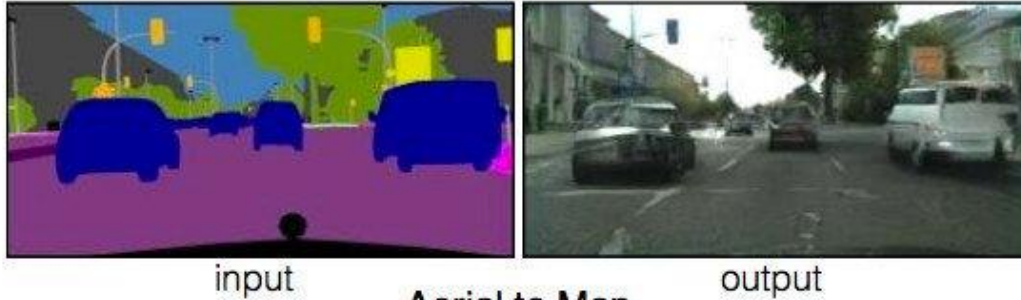
OUTPUT



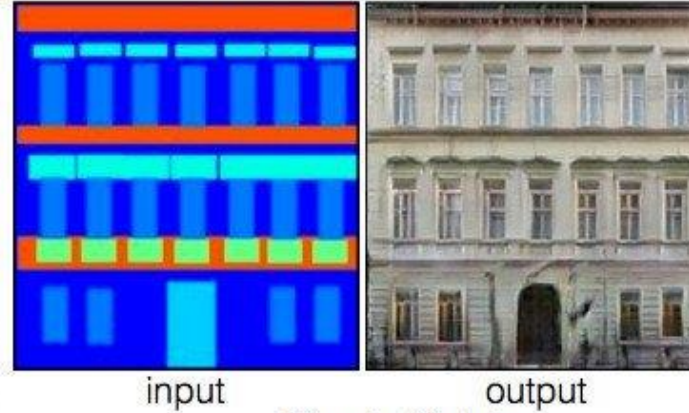
Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." arXiv preprint (2017).

# Image-to-image translation

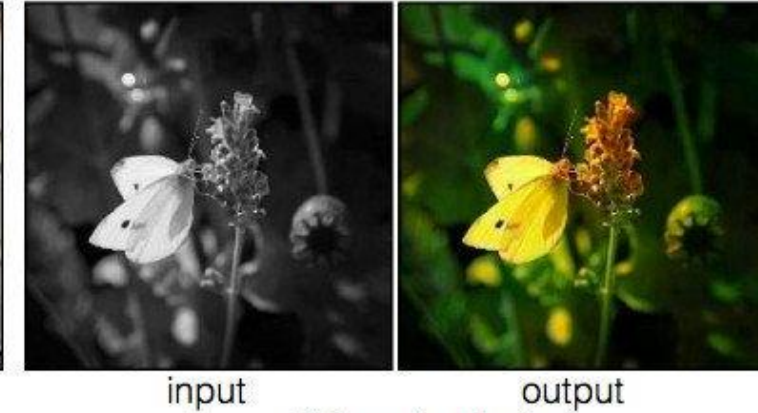
Labels to Street Scene



Labels to Facade



BW to Color



Aerial to Map



Day to Night



Edges to Photo



Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." arXiv preprint (2017).

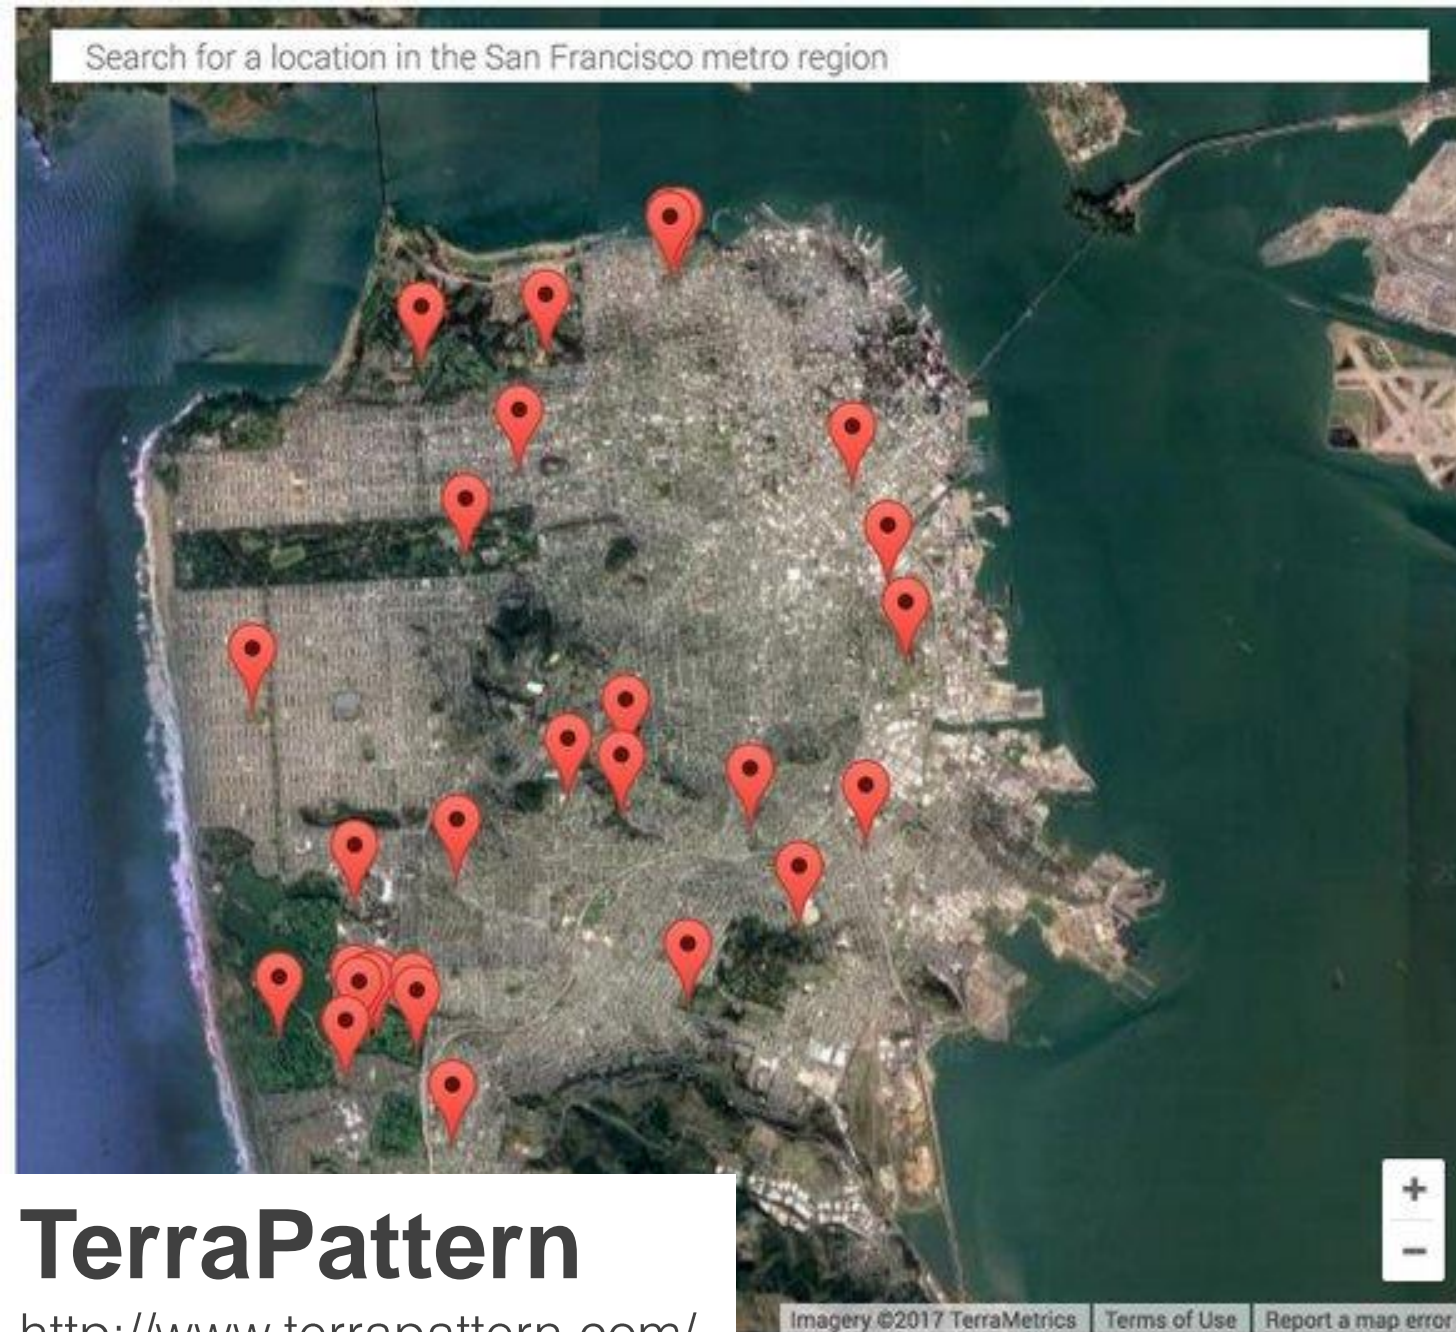


# Image-to-image translation



Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." arXiv preprint (2017).

Search for a location in the San Francisco metro region




# TerraPattern

<http://www.terrapattern.com/>


Kyle Bradbury

Neural Networks I

Geographical Plot



Similarity Plot



Search Results



Lecture 18

7

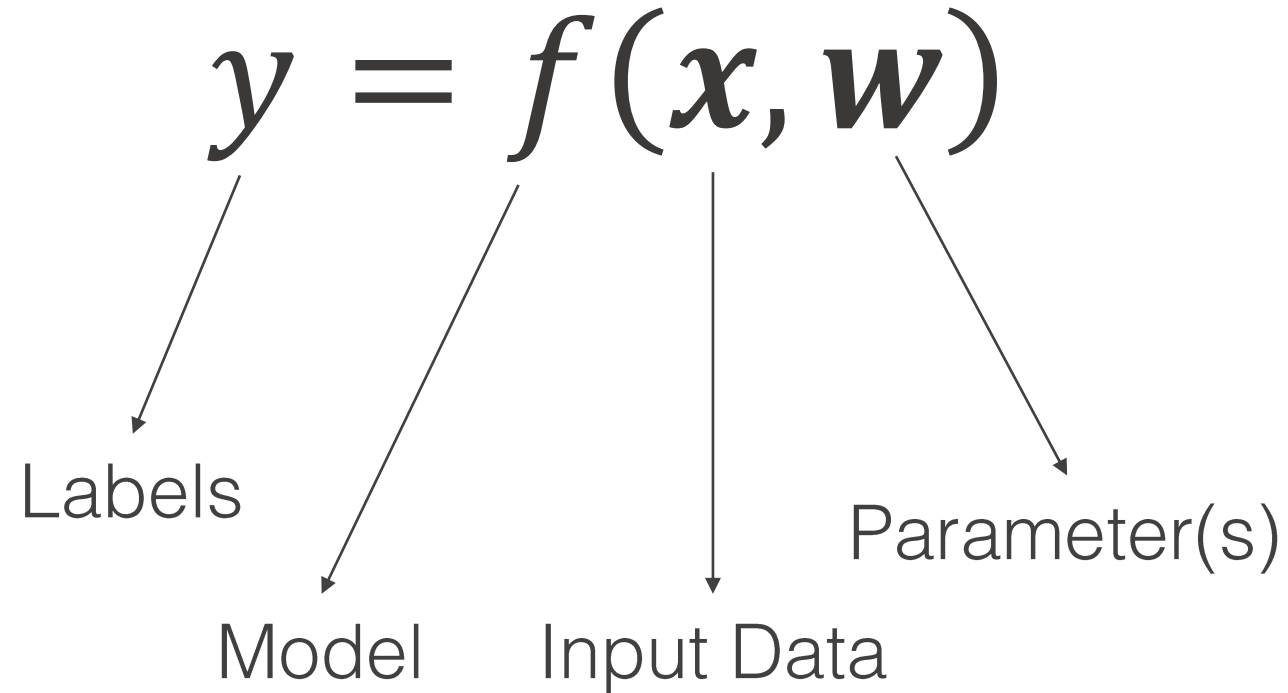
What is a neural network and **how does it work?**

How do we **choose model weights?**  
(i.e. how do we fit our model to data)

What are the challenges of using neural networks?




# Recall our goal in supervised learning



# Perceptron

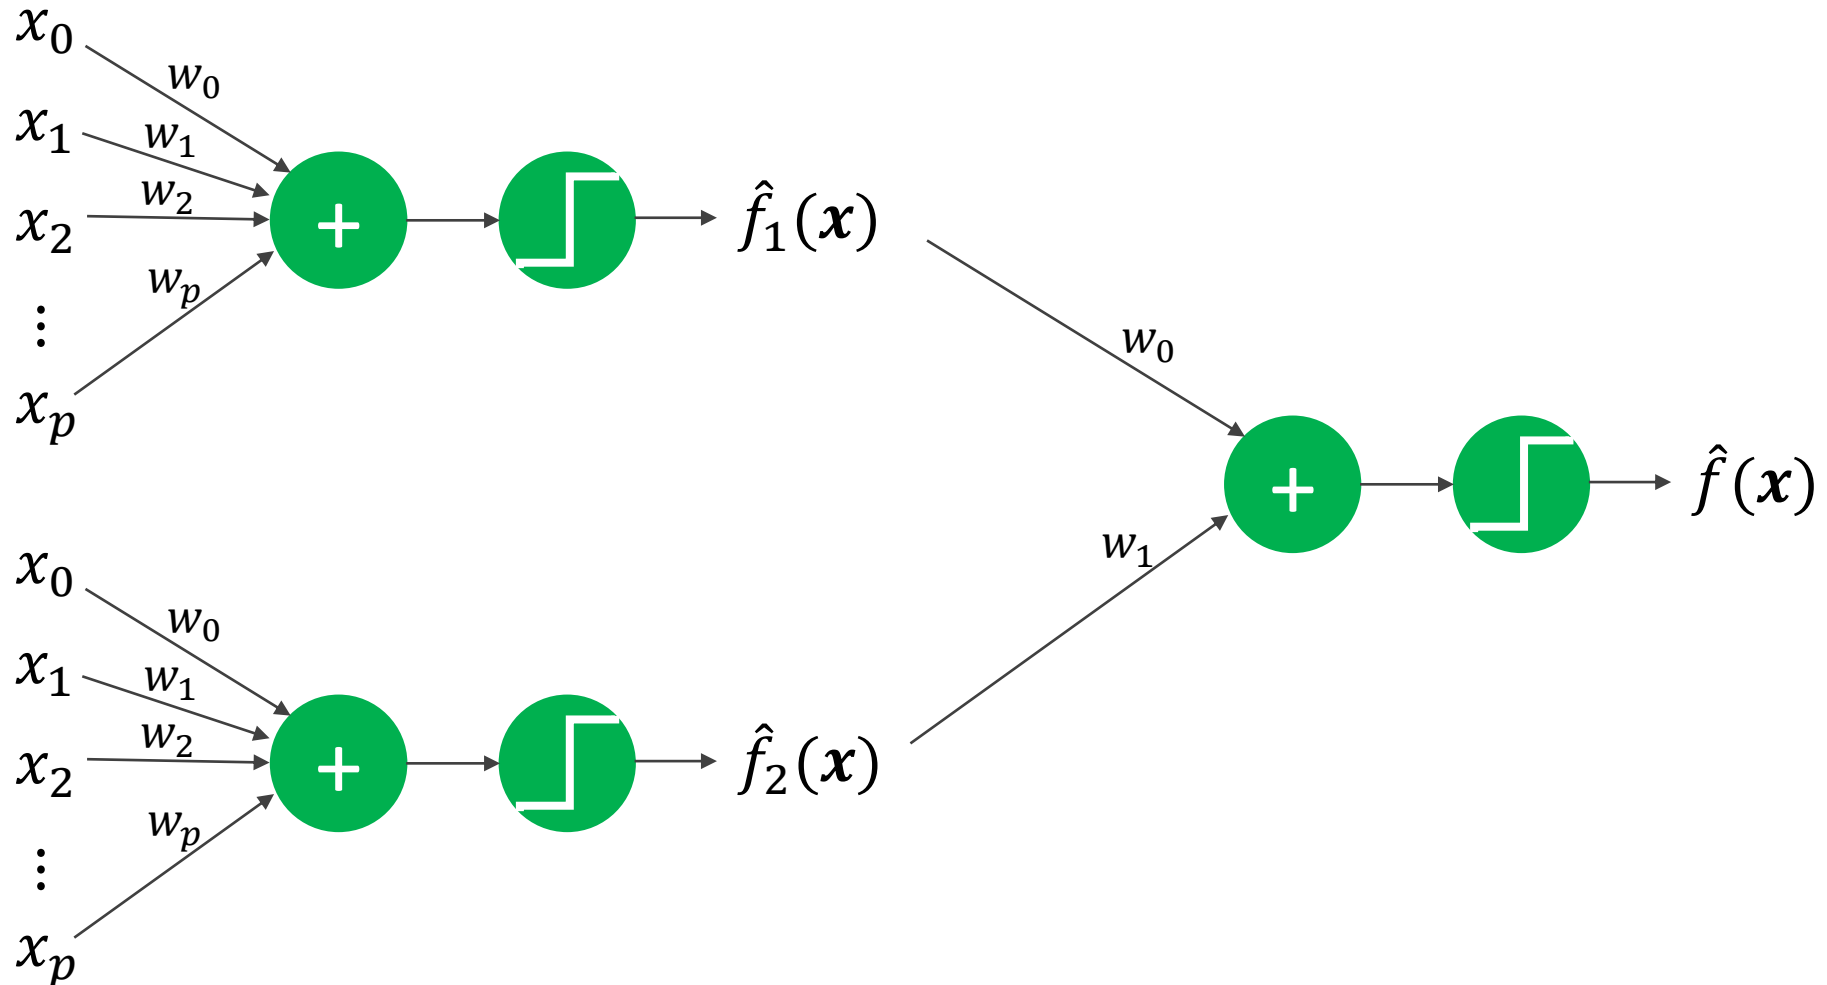
$$\hat{f}(x) = \text{sign} \left( \sum_{i=0}^p w_i x_i \right)$$



Source: Abu-Mostafa, Learning from Data, Caltech

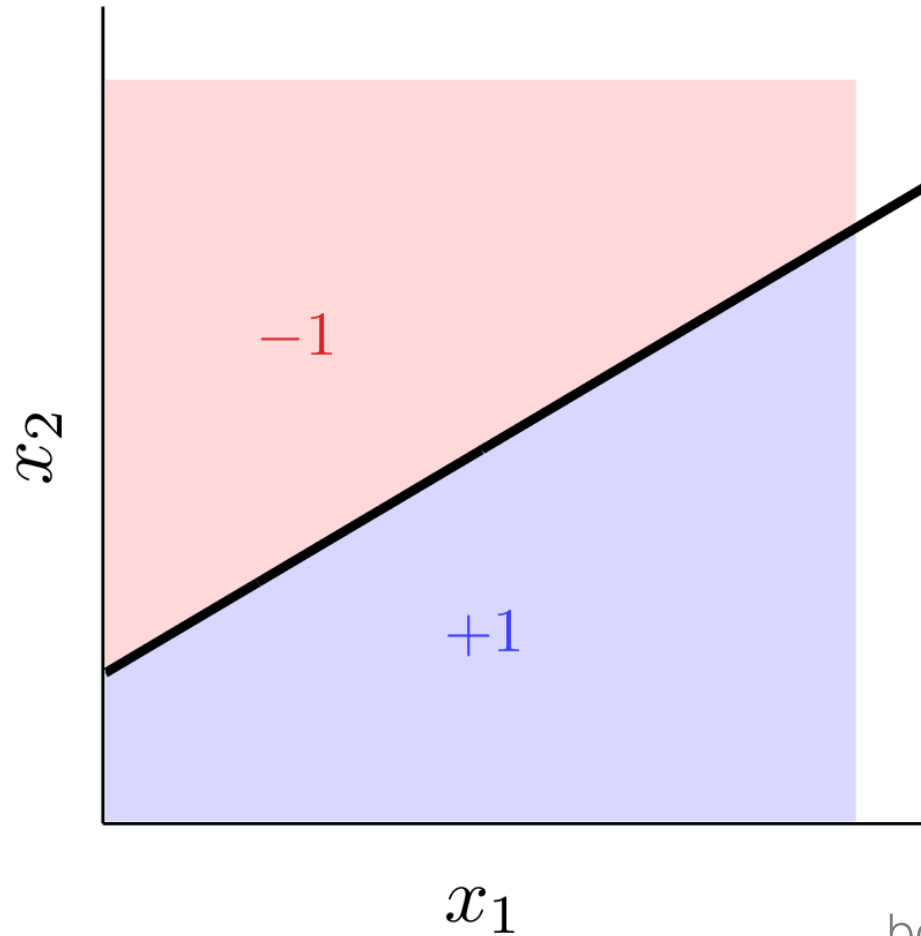
# Multilayer Perceptron

What if we stuck multiple perceptrons together?





## Perceptron #1

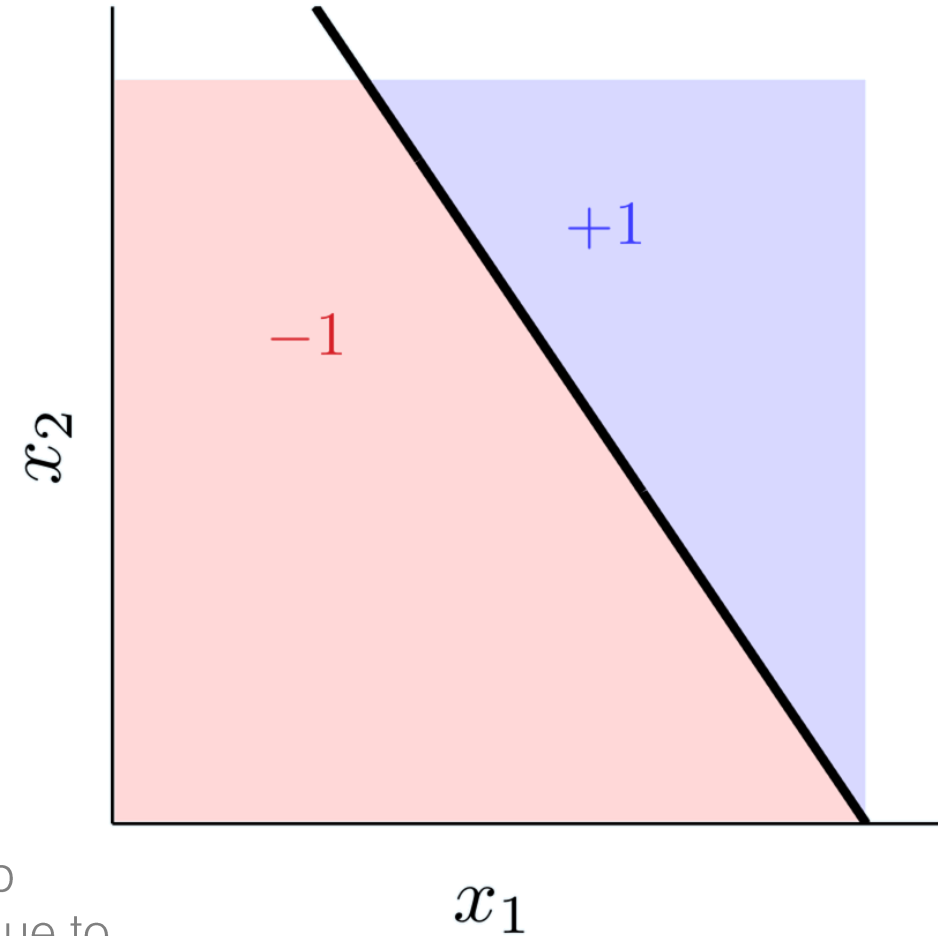


$$\hat{f}_1(\mathbf{x}) = \text{sign}(\mathbf{w}_1^T \mathbf{x})$$

The sharp  
boundary is due to  
our sign function



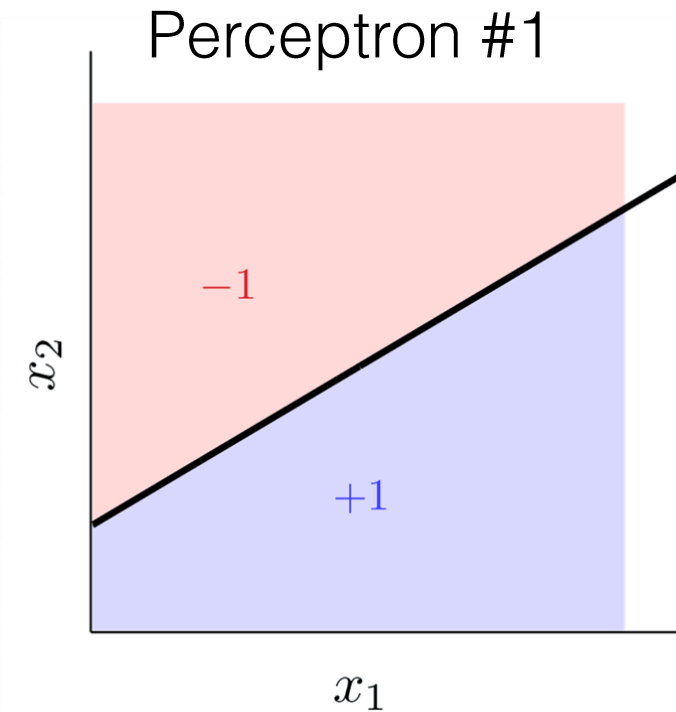
## Perceptron #2




$$\hat{f}_2(\mathbf{x}) = \text{sign}(\mathbf{w}_2^T \mathbf{x})$$

Source: Abu-Mostafa, Learning from Data, Caltech


**Multilayer perceptron:**  $\hat{f}(\mathbf{x}) = \begin{cases} +1 & \hat{f}_1(-\hat{f}_2) + (-\hat{f}_1)\hat{f}_2 > 0 \\ -1 & \text{else} \end{cases}$



$$\hat{f}_1(\mathbf{x}) = \text{sign}(\mathbf{w}_1^T \mathbf{x})$$



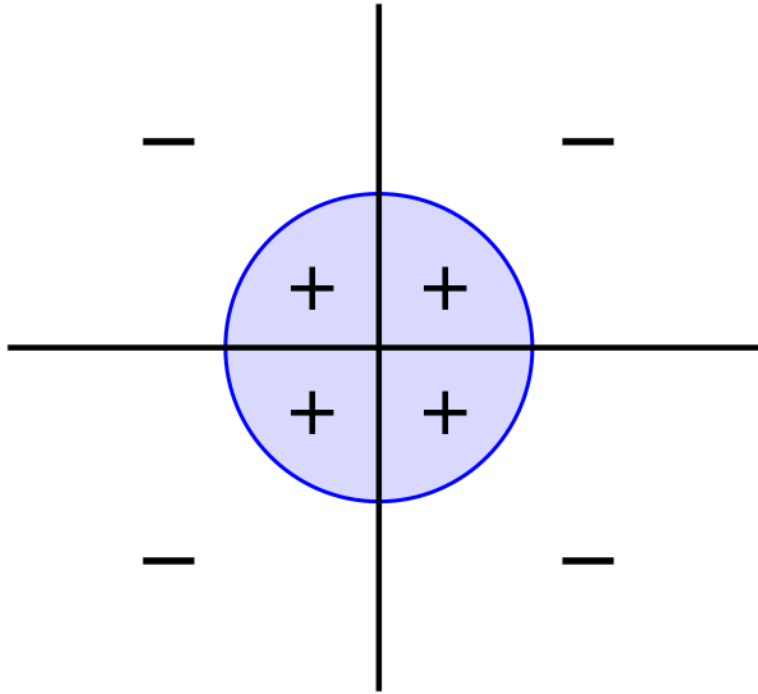
$$\hat{f}_2(\mathbf{x}) = \text{sign}(\mathbf{w}_2^T \mathbf{x})$$



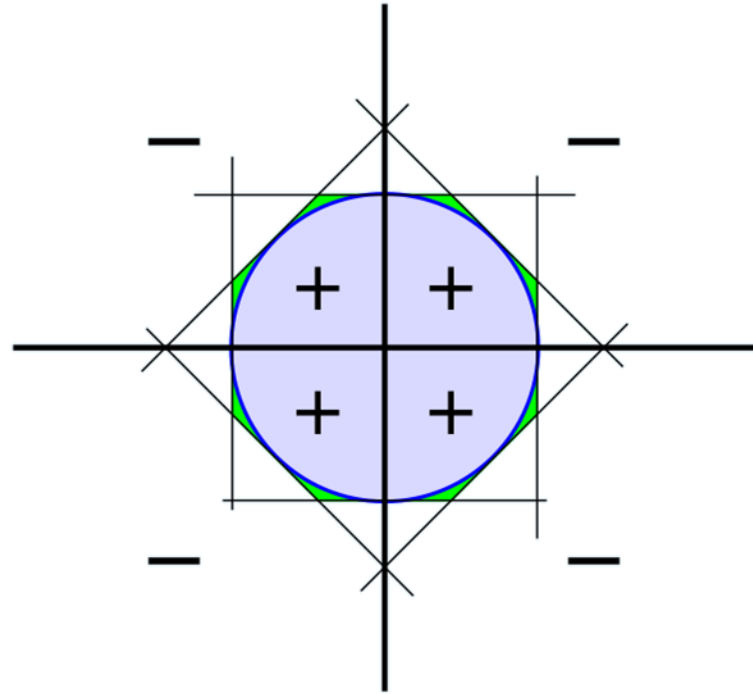
Source: Abu-Mostafa, Learning from Data, Caltech



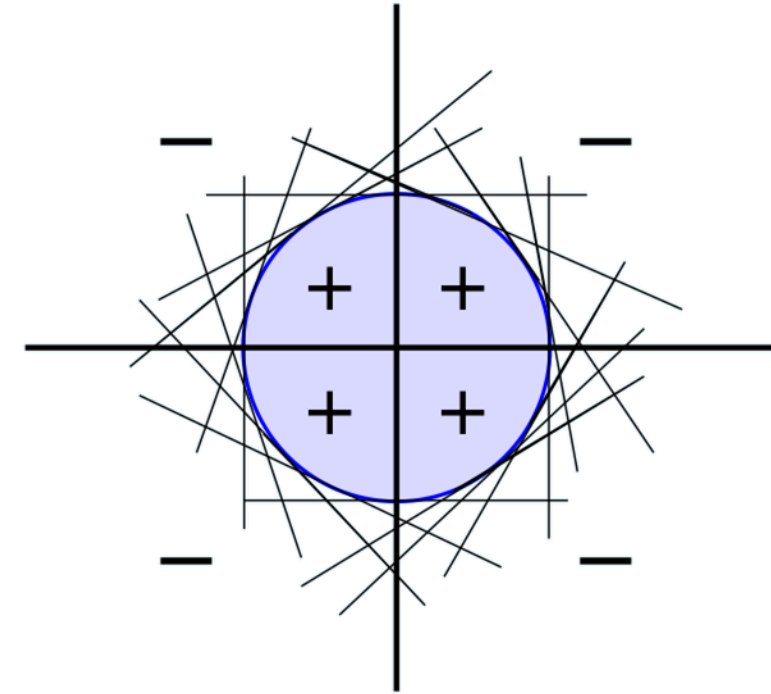
# Multilayer Perceptron



Target



8 perceptrons



16 perceptrons

The more nodes/neurons, the more flexible is the model

Source: Abu-Mostafa, Learning from Data, Caltech

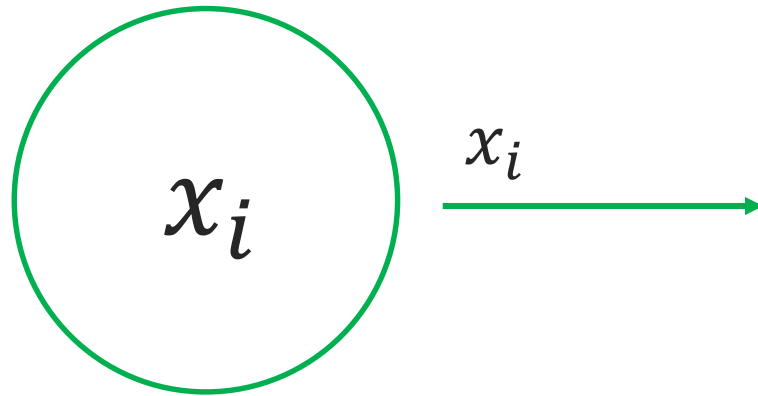
# Universal function approximation

“A **feedforward network** with a single layer is sufficient to represent **any function**, but the layer may be infeasibly large and may fail to learn and generalize correctly.”

Ian Goodfellow, Deep Learning  
Creator of generative adversarial networks



# Input nodes / neurons



Simply passes the input value to the next layer

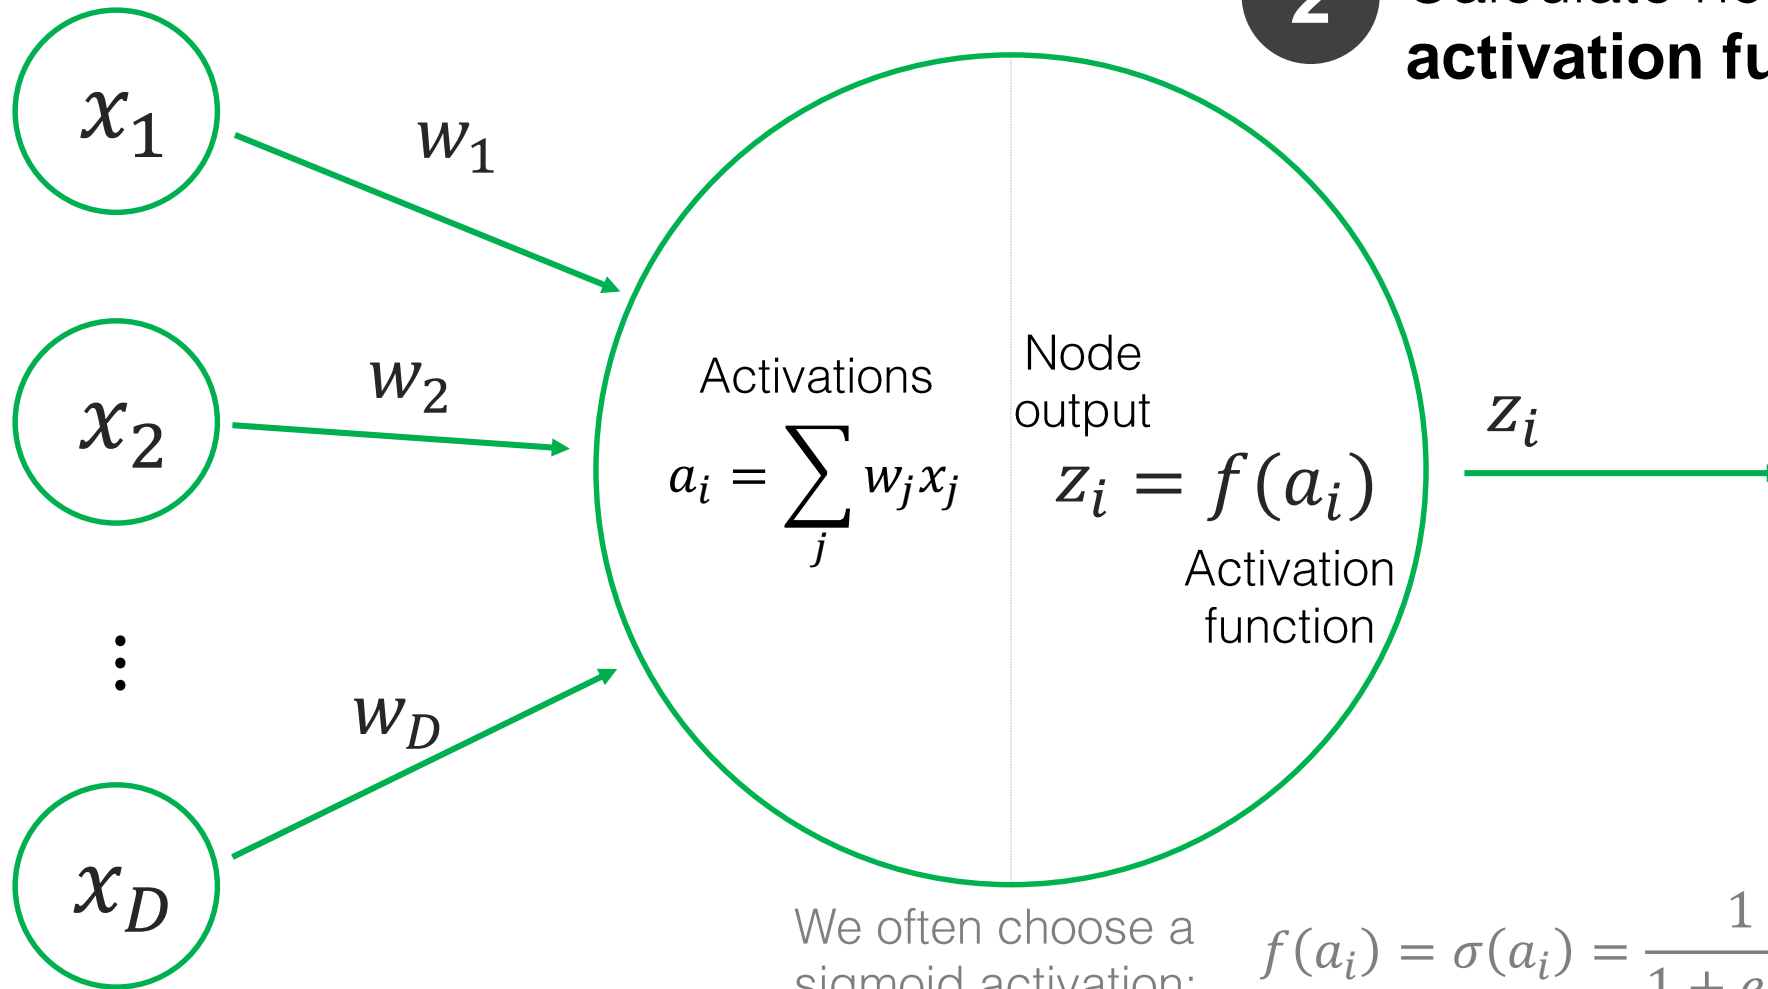
# Hidden & output nodes

1

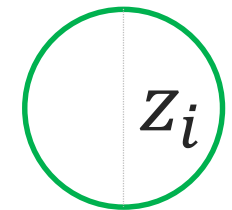
Calculate the **activations**: linear combinations of weights and the last layer's output

2


Calculate node output: apply the **activation function** to the activations



Represented as:



# Simple Neural Network



**Notational shorthand:**  
(a more precise alternative notation)

$$w_{ij} = w_{ij}^{(1)}$$

$$z_i = z_i^{(1)}$$

$$w_{ij} = w_{ij}^{(2)}$$

$$y = z_1^{(2)}$$

$w_{ij}^{(k)}$

- Layer  $k$
- From node  $j$  (in the last layer)
- to node  $i$  (in the next layer)



# Forward Propagation


Calculating the output from input

$$a_1 = (0.9)(1.0) + (0.3)(0.5) = 1.05 \quad \text{Hidden layer calculations}$$

$$a_2 = (0.2)(1.0) + (0.8)(0.5) = 0.6$$

$$z_1 = \sigma(a_1) = \sigma(1.05) = 0.741$$

$$z_2 = \sigma(a_2) = \sigma(0.6) = 0.646$$



Output layer calculations


$$a_y = (0.2)(0.741) + (0.5)(0.646) = 0.471$$

$$y = \sigma(a_y) = \sigma(0.471) = 0.616$$

$$\sigma(a_i) = \frac{1}{1 + e^{-a_i}}$$

# Forward Propagation

Calculating the output from input



## Hidden layer matrix calculations

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \end{bmatrix} \quad \mathbf{z} = \begin{bmatrix} z_1 \\ z_2 \end{bmatrix}$$

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} \begin{array}{l} \longrightarrow \text{The weights INTO node } z_1 \\ \longrightarrow \text{The weights INTO node } z_2 \end{array}$$

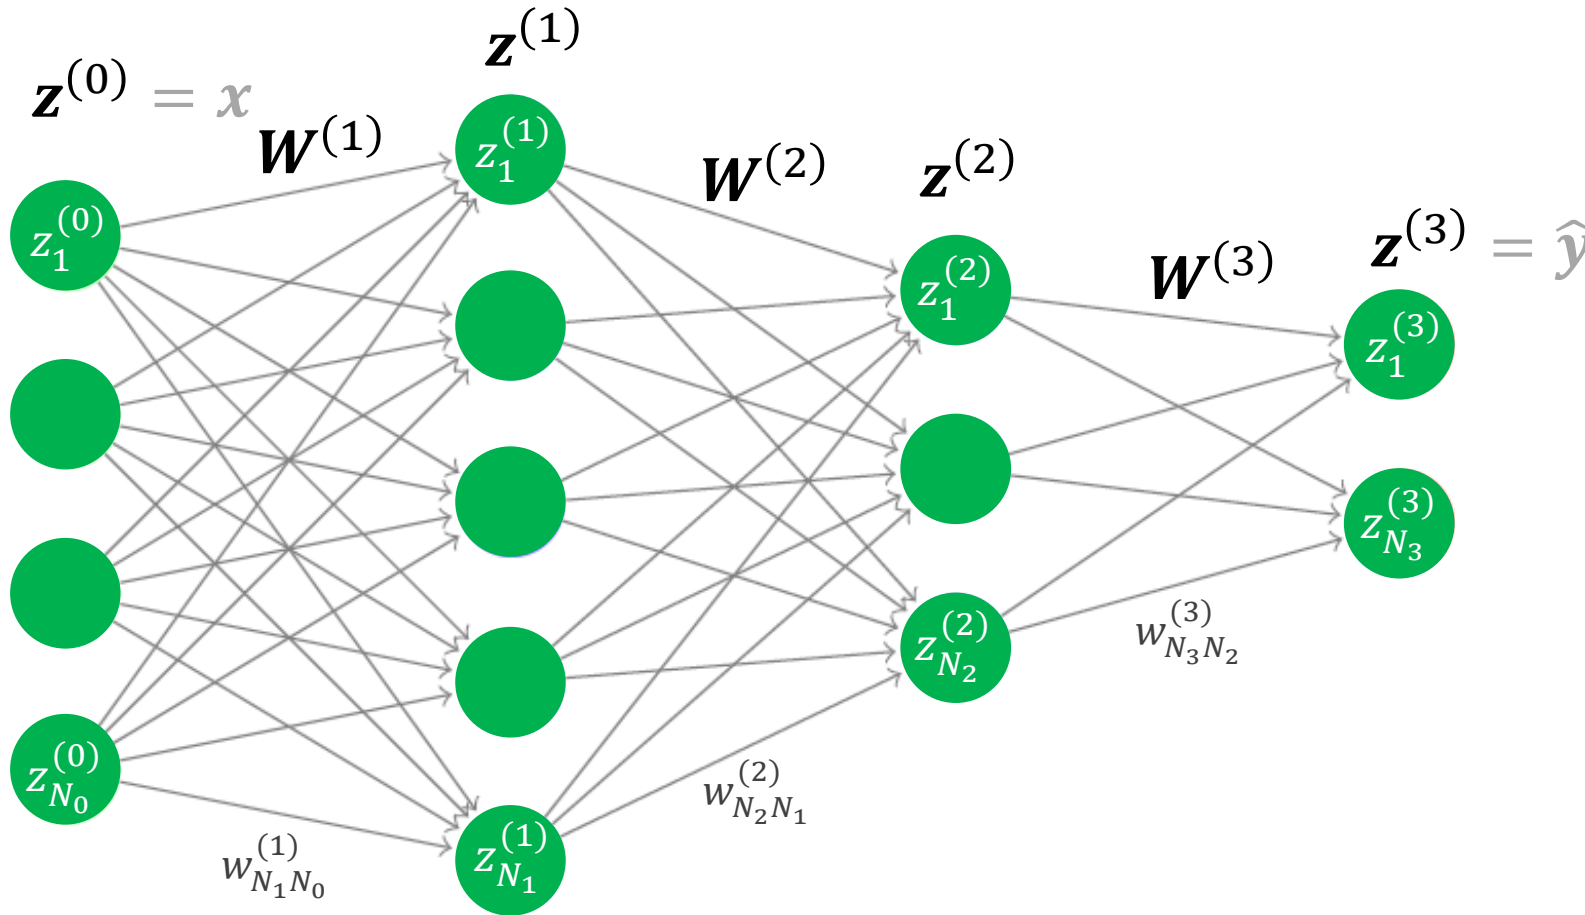
$$\mathbf{a} = \mathbf{W}\mathbf{x} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

$$= \begin{bmatrix} w_{11}x_1 + w_{12}x_2 \\ w_{21}x_1 + w_{22}x_2 \end{bmatrix}$$

$$\mathbf{z} = \sigma(\mathbf{a}) = \begin{bmatrix} \sigma(w_{11}x_1 + w_{12}x_2) \\ \sigma(w_{21}x_1 + w_{22}x_2) \end{bmatrix}$$

# Forward Propagation

Example neural network with  $L = 3$  layers and the  $i$ th layer has  $N_i$  nodes



Simple steps for forward propagation:

$$\text{For } i = 1 \text{ to } L - 1:$$
$$\mathbf{z}^{(i)} = \sigma(\mathbf{W}^{(i)} \mathbf{z}^{(i-1)})$$

Where:

$$\mathbf{z}^{(0)} = \mathbf{x}$$


$$\hat{\mathbf{y}} = \mathbf{z}^{(L)}$$

Prediction error is measured:

$$E_n = \frac{1}{2} (\hat{y}_n - y_n)^2$$



# Neural networks can be customized














# Input

# Hidden 1


# Hidden 2

# Output




Multilayer neural nets can build up from basic building blocks to more complex structures

# From regression to classification



For **binary classification** with a sigmoid activation function, the output is between zero and one, so threshold this value to assign the class

# From regression to classification



For **multiclass problems**, we can have multiple outputs and use a softmax function:

$$y_i = g(a_i) = \frac{e^{a_i}}{\sum_{n=1}^M e^{a_n}}$$

Choose the largest  $y$  value as the predicted class

As with many aspects of neural networks this is one of a number of approaches



# Next time...

What is a neural network and **how does it work?**

How do we **choose model weights?**

(i.e. how do we fit our model to data)

What are the challenges of using neural networks?